

# Outline

❖ Introduction

❖ Ensemble Learning

❖ Boosting Algorithms

❖ Demo in Python

# Part 1:

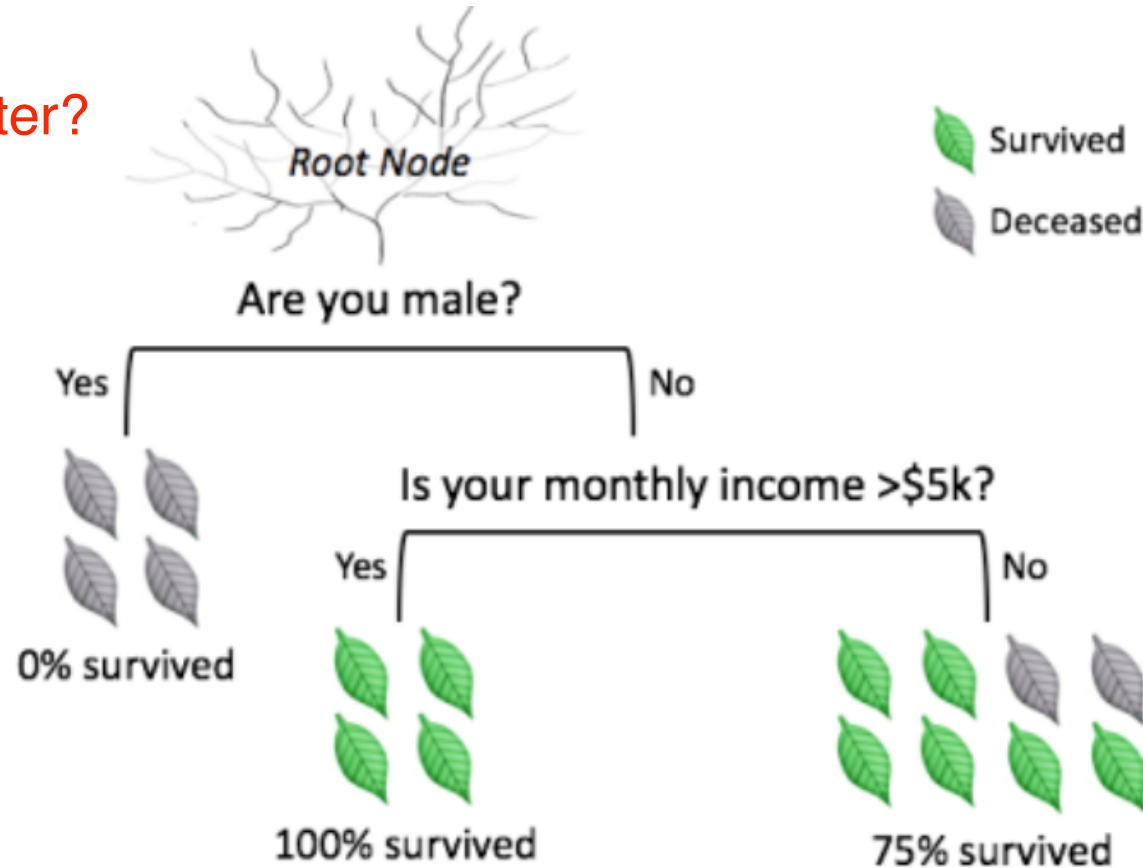
## Overview of Decision Trees

# Decision Trees

Would you survive a disaster?

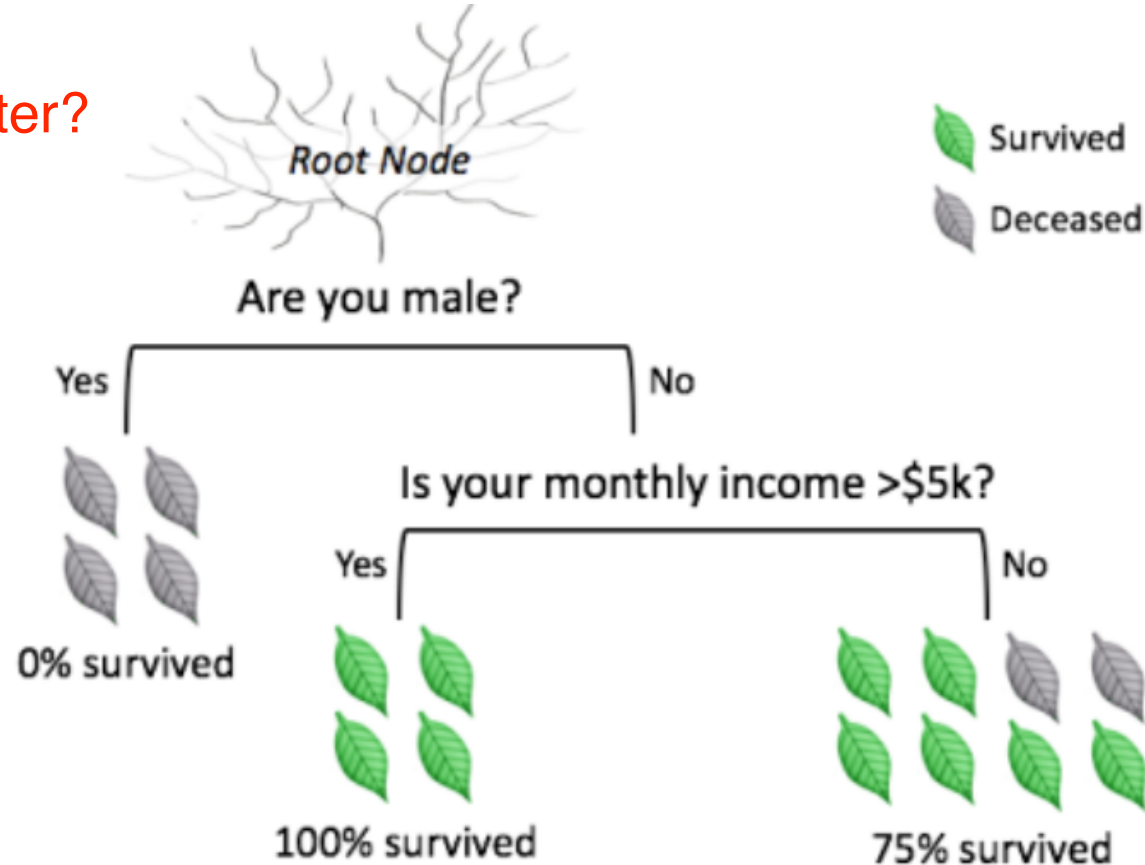
# Decision Trees

Would you survive a disaster?



# Decision Trees

Would you survive a disaster?

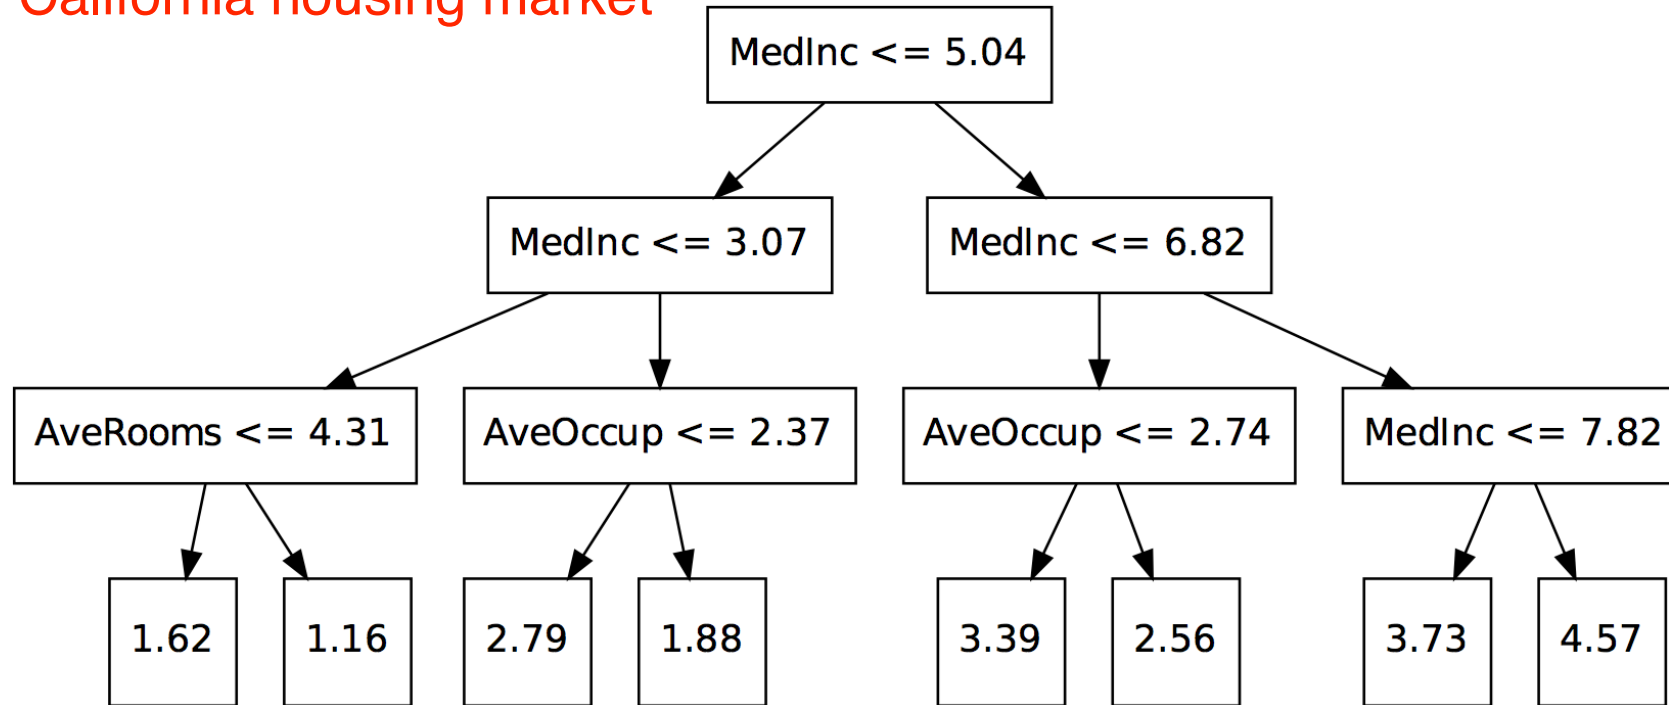


## Stopping criteria:

- Stop when data points at the leaf are all of the same class
- Stop when the leaf contains less than K data points
- Stop when further branching does not improve homogeneity beyond a minimum threshold

# Decision Tree Overview

Predictions for California housing market



# Decision Trees: Practical Use

## Strengths

- Non linear
- Robust to correlated features
- Robust to feature distributions
- Robust to missing values
- Simple to comprehend
- Fast to train
- Fast to score

## Weaknesses

- Poor accuracy
- Cannot project
- Inefficiently fits linear relationships

## Part 2:

# Overview of Ensemble Learning

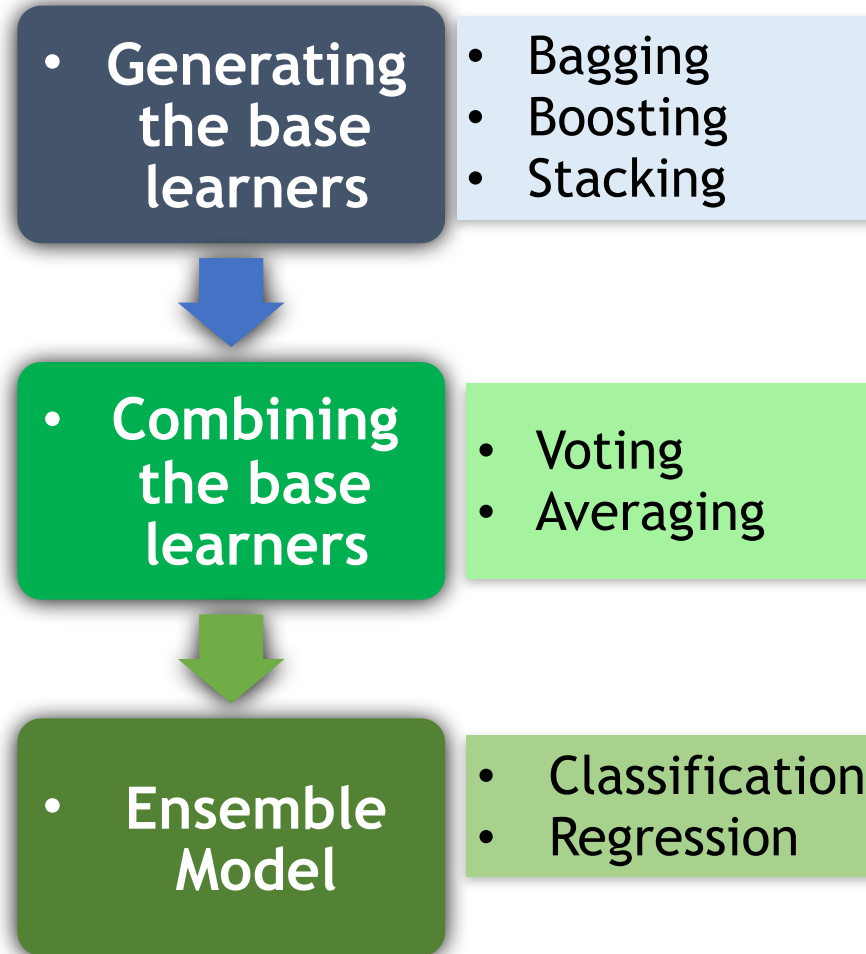


# Ensemble Learning

A **Mixture** of experts or “base learners” combined to solve the same learning problem.

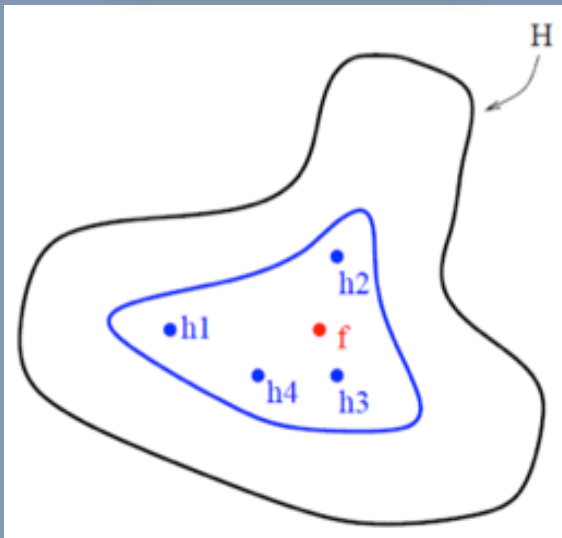
Construct a **set of hypotheses** and combine them to use .

The **generalization** ability is much stronger than that of base learners.



# Why Ensembles Superior to Singles?

## Statistical



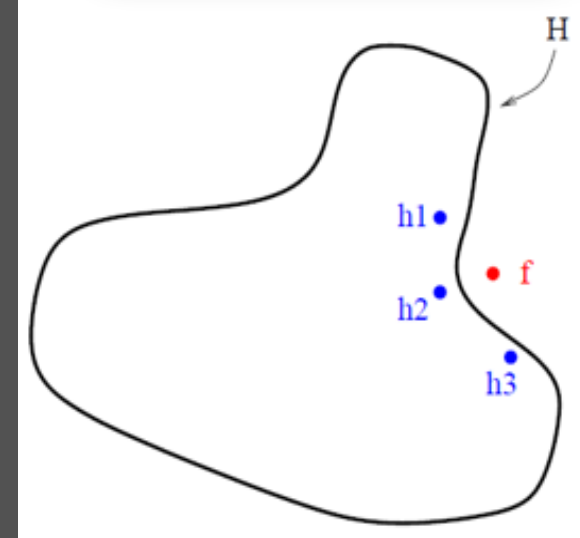
No sufficient training to choose a single best learner.

## Computational



Imperfect search processes result in sub-optimal hypotheses.

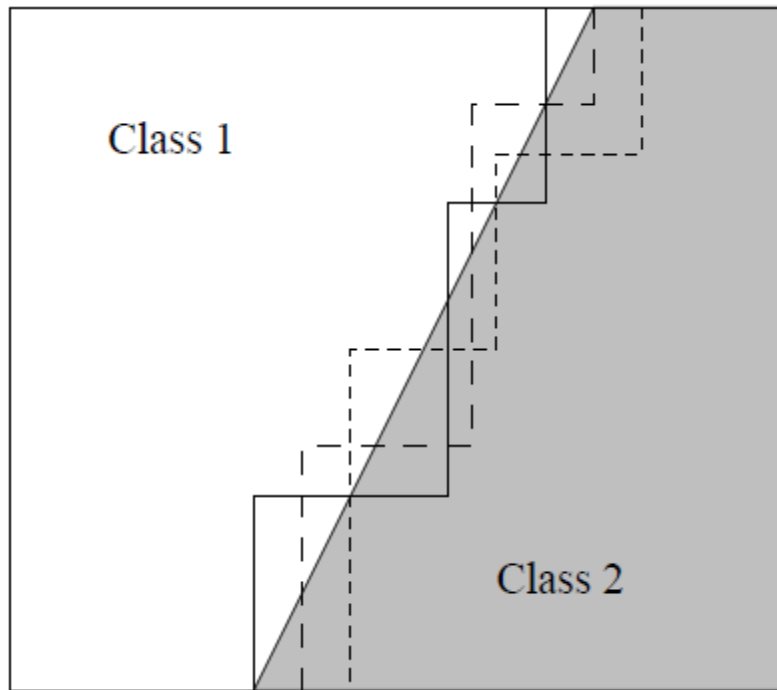
## Representational



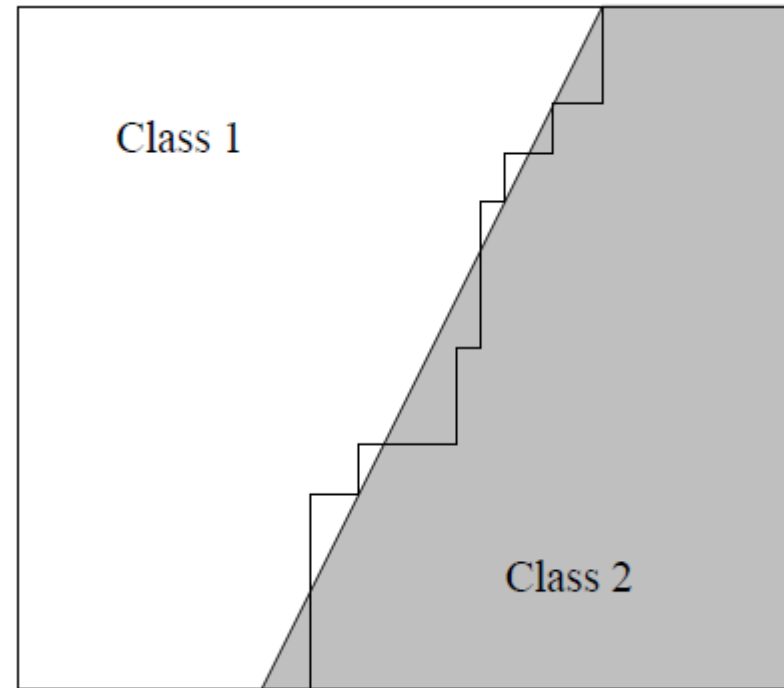
The true function cannot be represented by any of the hypothesis in  $H$ .

# Illustration of the representation issue

**Diagonal Decision boundary with decision trees base learners**

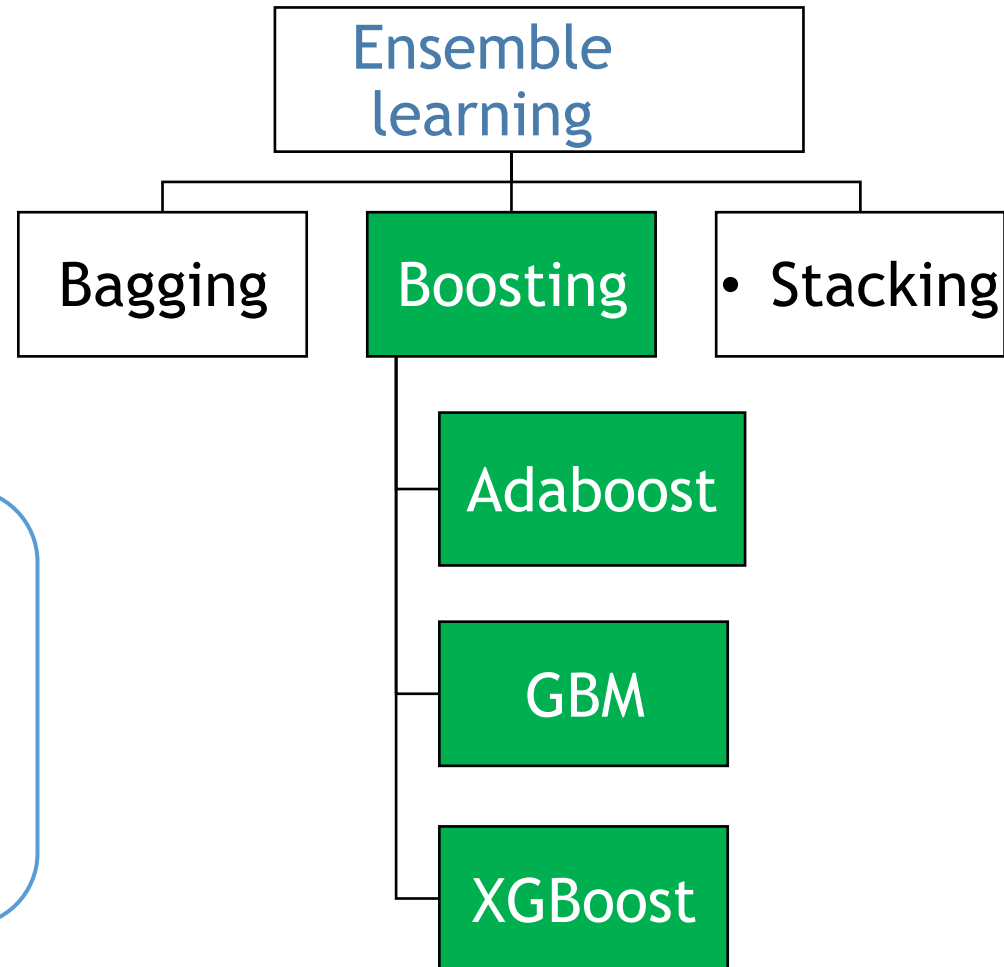


Three staircase approximations



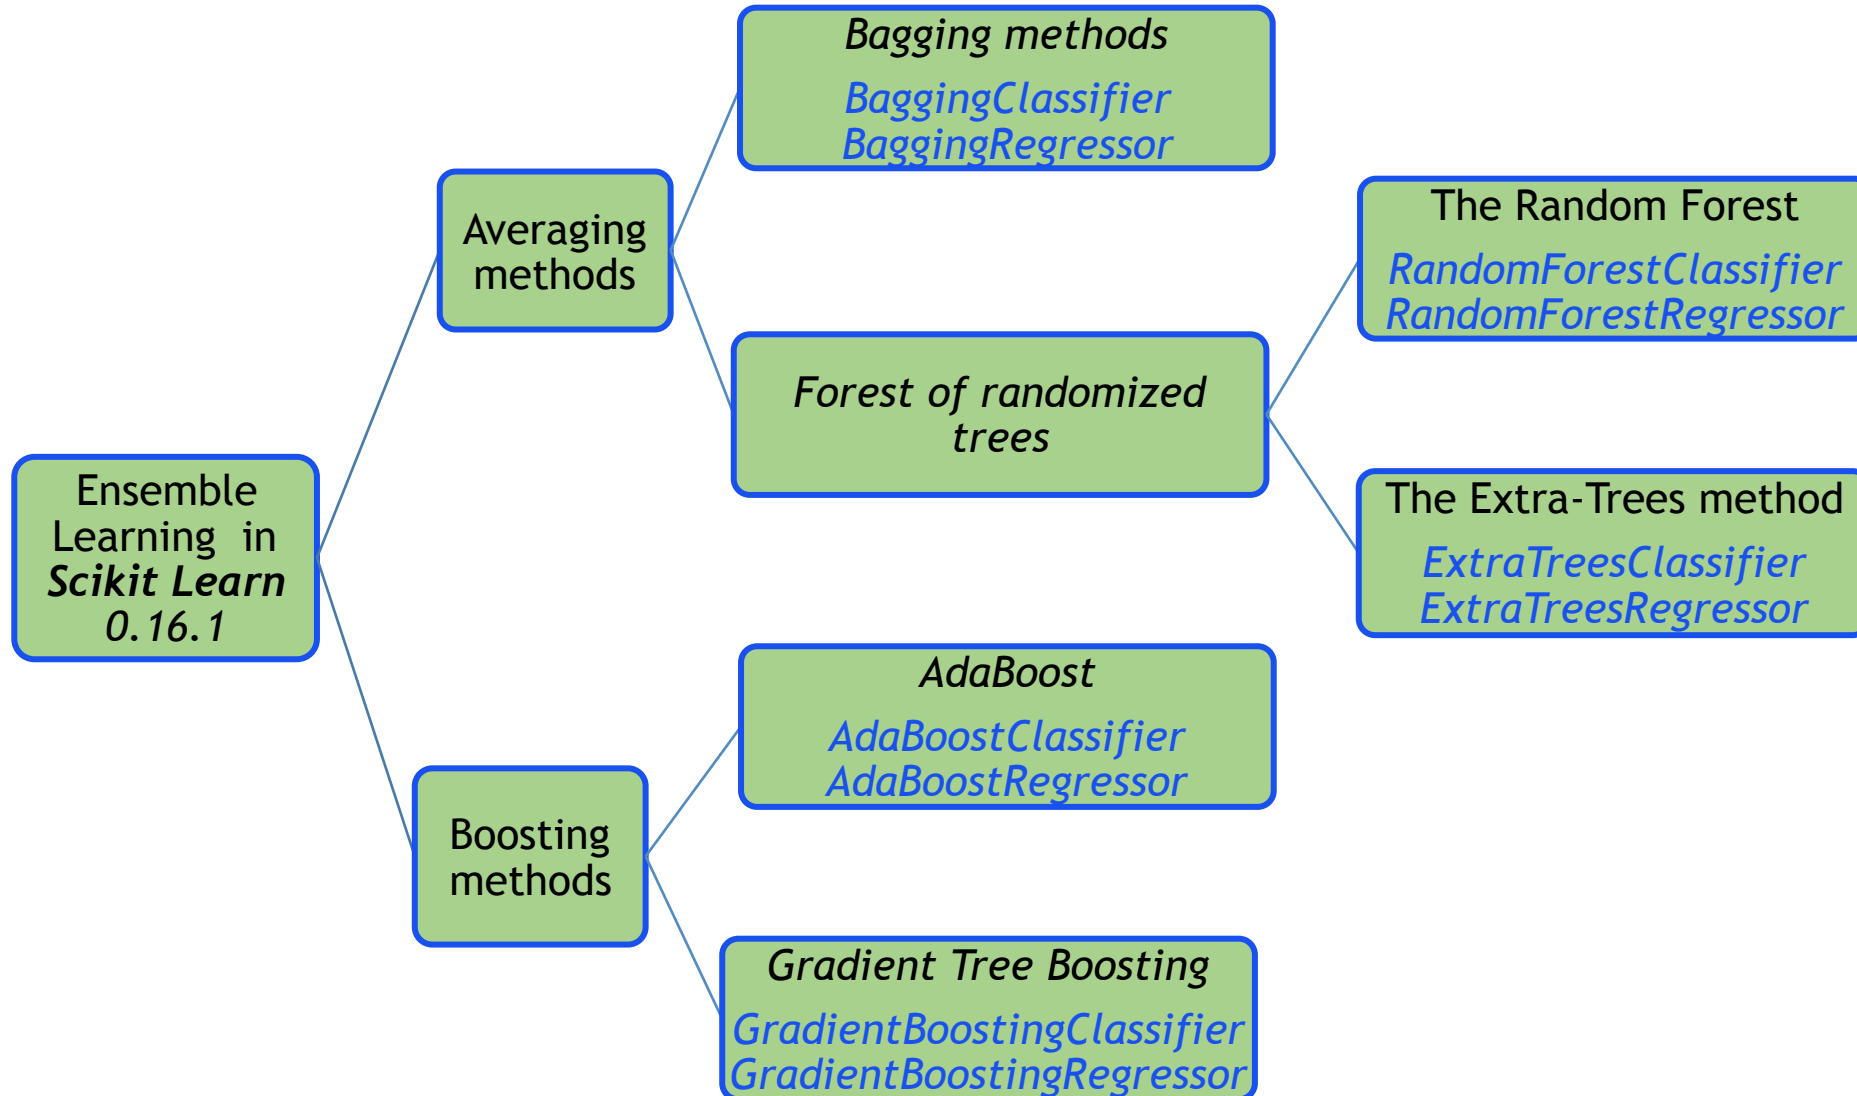
The voted decision boundary

# Boosting is one Type of Ensemble Learning



- Bagging & Boosting are considered homogeneous. They use a single base learning algorithm.
- Staking uses different kinds of base learners

# Ensemble Learning in SKlearn



# Bagging

- Trains a number of base learners each from a different *bootstrap sample* (**Bootstrap AGG**regat**ING**)
- Each dataset is generated by sampling from the total N data examples, choosing N items uniformly at random **with replacement**.
- For a bootstrap sample, some training examples may appear but some may not.
- The outputs of the models are combined by:
  - Averaging (in the case of regression)
  - Voting (in the case of classification)
- Example: Random forests

# Bagging

---

## Algorithm    Bagging

---

**Input:** Required ensemble size  $T$

**Input:** Training set  $S = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$

**for**  $t = 1$  to  $T$  **do**

    Build a dataset  $S_t$ , by sampling  $N$  items, randomly *with replacement*

    Train a model  $h_t$  using  $S_t$ , and add it to the ensemble.

**end for**

For a new testing point  $(x', y')$ ,

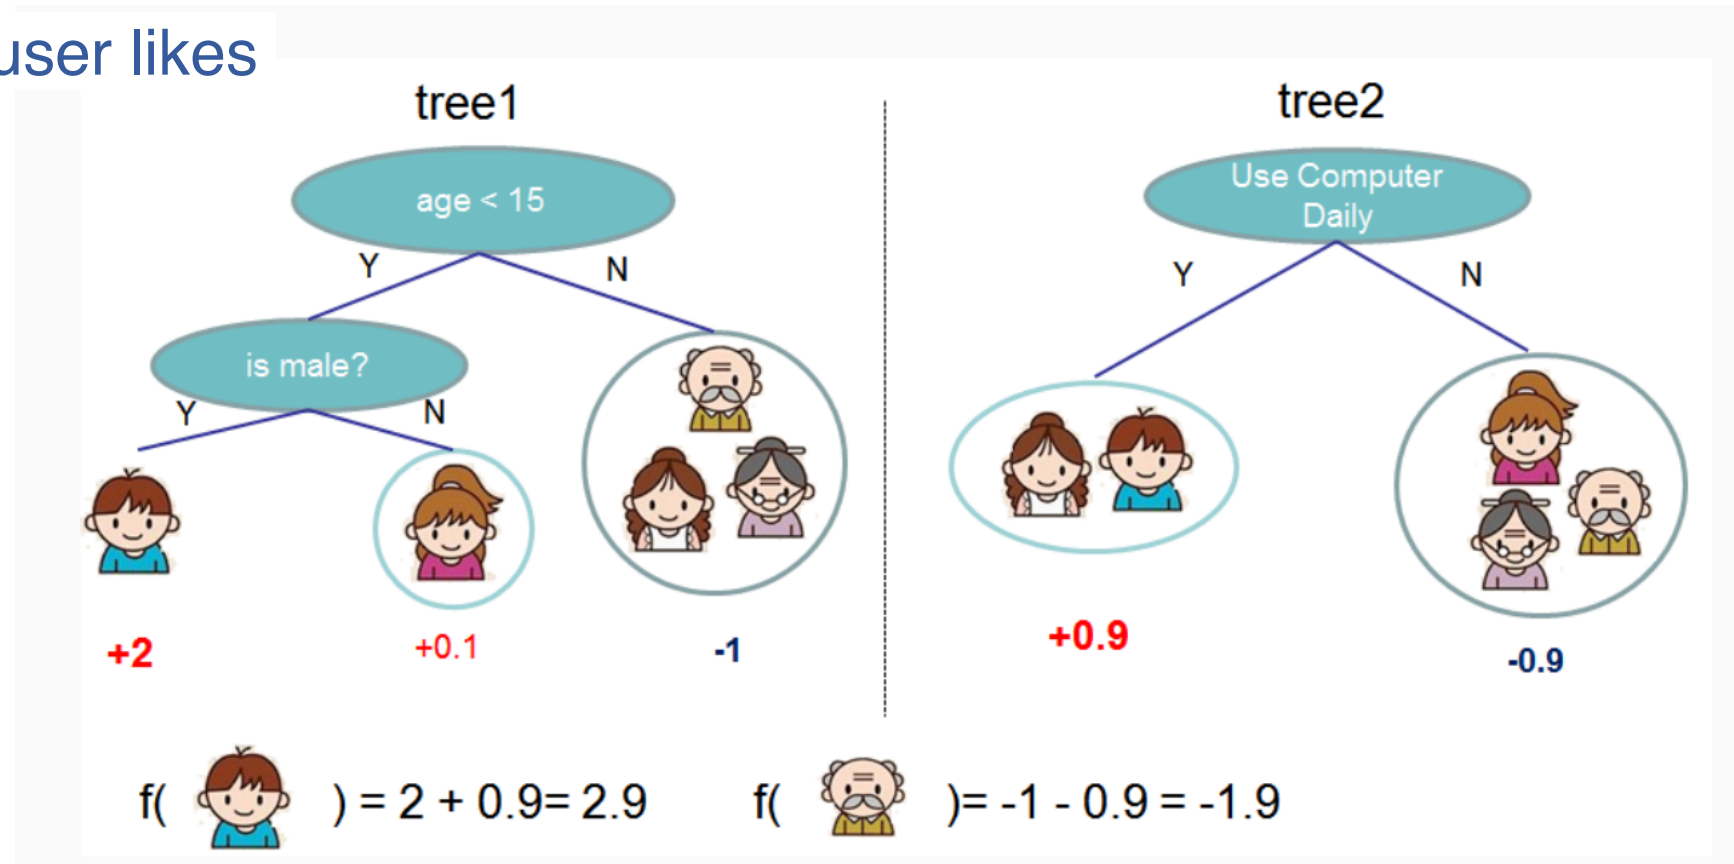
If model outputs are continuous, combine them by averaging.

If model outputs are class labels, combine them by voting.

---

# Tree Ensemble Model

Predict whether a given user likes computer games or not





# Decision Trees & Random forests in SKlearn

`sklearn.tree.DecisionTreeClassifier`

Input parameters:

`criterion`  
`max_depth`  
`min_samples_split`  
`min_samples_leaf`  
`max_features`  
`max_leaf_nodes`  
`...`

`sklearn.ensemble.RandomForestClassifier`

Input parameters:

`n_estimators`  
`+`  
`criterion`  
`max_depth`  
`min_samples_split`  
`min_samples_leaf`  
`max_features`  
`max_leaf_nodes`  
`...`

# Variants of Bagging

## Pasting

When random subsets of the dataset are drawn as random subsets of the samples.

L. Breiman, "Pasting small votes for classification in large databases and on-line", Machine Learning, 36(1), 85-103, 1999.

## Bagging

When random samples of the dataset are drawn with replacement

L. Breiman, "Bagging predictors", Machine Learning, 24(2), 123-140, 1996

## Random Subspaces

When random subsets of the dataset are drawn as random subsets of the features

•T. Ho, "The random subspace method for constructing decision forests", Pattern Analysis and Machine Intelligence, 20(8), 1998.

## Random Patches

When base estimators are built on subsets of both samples and features

G. Louppe and P. Geurts, "Ensembles on Random Patches", Machine Learning and Knowledge Discovery in Databases 2012

## Random Forests

A hybrid of the Bagging and the Random Subspace Method Uses Decision Trees as the base classifier with random splits

L. Breiman, "Random Forests", Machine Learning, 45(1), 5-32, 2001.

## Random Forests

- The scikit-learn implements by averaging instead of voting.
- The split that is picked is the best split among a random subset of the features.

## Extremely Randomized Trees

- Randomness goes one step further in the way splits are computed.
- As in random forests, a random subset of candidate features is used.
- Thresholds are drawn at random for each candidate feature and the best of these randomly-generated thresholds is picked as the splitting rule.

Main  
parameter

number of trees in the forest.



“n\_estimators”

Feature  
importance

1. Top of the tree. 2. Used in many trees



Feature  
Selection

# Stacking

**Input:** Data set  $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\}$ ;

First-level learning algorithms  $\mathcal{L}_1, \dots, \mathcal{L}_T$ ;

Second-level learning algorithm  $\mathcal{L}$ .

**Process:**

for  $t = 1, \dots, T$ :

$h_t = \mathcal{L}_t(\mathcal{D})$       % Train a first-level individual learner  $h_t$  by applying the first-level

end;      % learning algorithm  $\mathcal{L}_t$  to the original data set  $\mathcal{D}$

$\mathcal{D}' = \emptyset$ ;      % Generate a new data set

for  $i = 1, \dots, m$ :

for  $t = 1, \dots, T$ :

$z_{it} = h_t(\mathbf{x}_i)$       % Use  $h_t$  to classify the training example  $\mathbf{x}_i$

end;

$\mathcal{D}' = \mathcal{D}' \cup \{((z_{i1}, z_{i2}, \dots, z_{iT}), y_i)\}$

end;

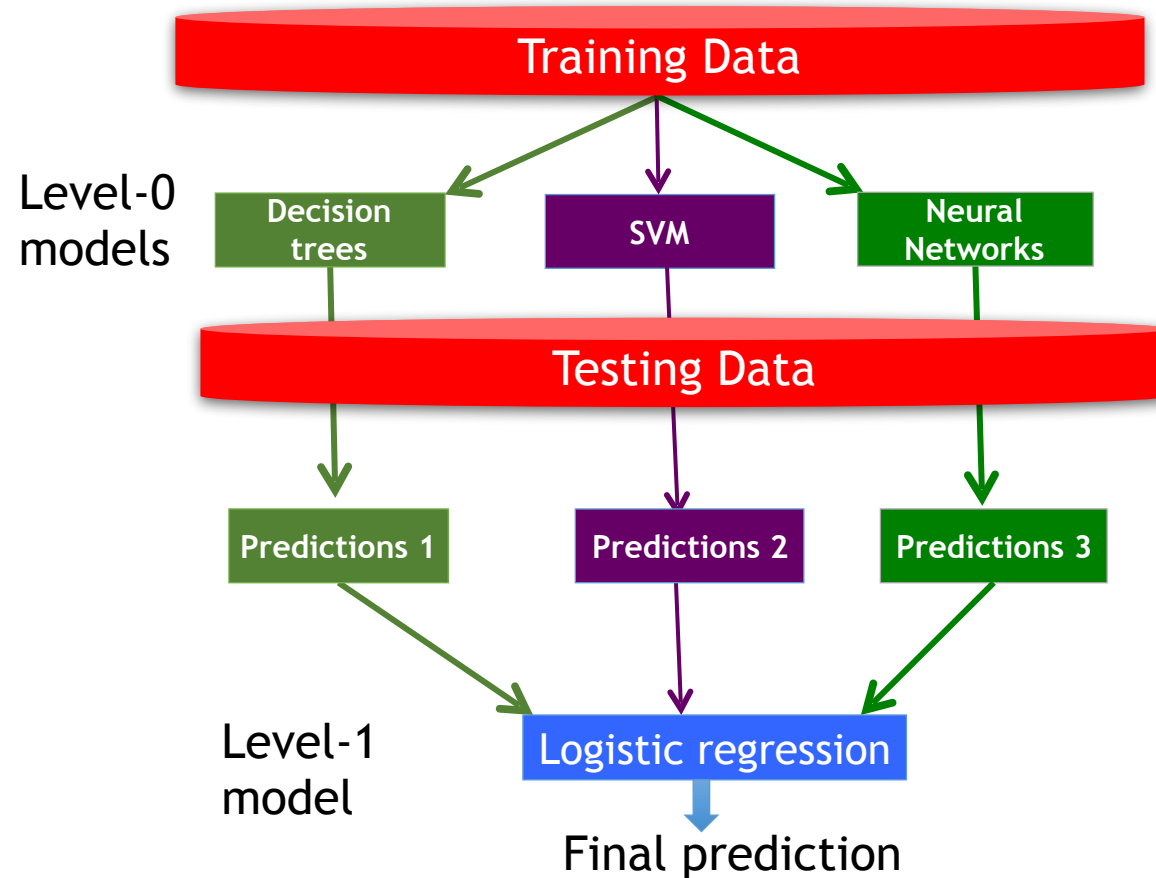
$h' = \mathcal{L}(\mathcal{D}')$ .      % Train the second-level learner  $h'$  by applying the second-level

% learning algorithm  $\mathcal{L}$  to the new data set  $\mathcal{D}'$

**Output:**  $H(\mathbf{x}) = h' (h_1 (\mathbf{x}), \dots, h_T (\mathbf{x}))$

# Stacking

Stacked generalization (or stacking) is used to combine models of different types.



# Part 3:

## Boosting Algorithms

# AdaBoost

- It Creates a ‘weak’ classifier that its accuracy is only slightly better than random guessing.
- A succession of models are built iteratively.
- **Records that were misclassified by the previous model are given more weight.**
- Finally, all of the successive models are weighted according to their success.
- Uses decision stumps as the base learners

Schapire, R.E.: The strength of weak learnability. Machine Learning 5(2) (1990)

# AdaBoost

---

**Algorithm** Adaboost

---

**Input:** Required ensemble size  $T$

**Input:** Training set  $S = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$ , where  $y_i \in \{-1, +1\}$   
Define a uniform distribution  $D_1(i)$  over elements of  $S$ .


**for**  $t = 1$  to  $T$  **do**

Train a model  $h_t$  using distribution  $D_t$ .

Calculate  $\epsilon_t = P_{D_t}(h_t(x) \neq y)$

If  $\epsilon_t \geq 0.5$  break

Set  $\alpha_t = \frac{1}{2} \ln \left( \frac{1-\epsilon_t}{\epsilon_t} \right)$

Update  $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$  

where  $Z_t$  is a normalization factor so that  $D_{t+1}$  is a valid distribution.

**end for**

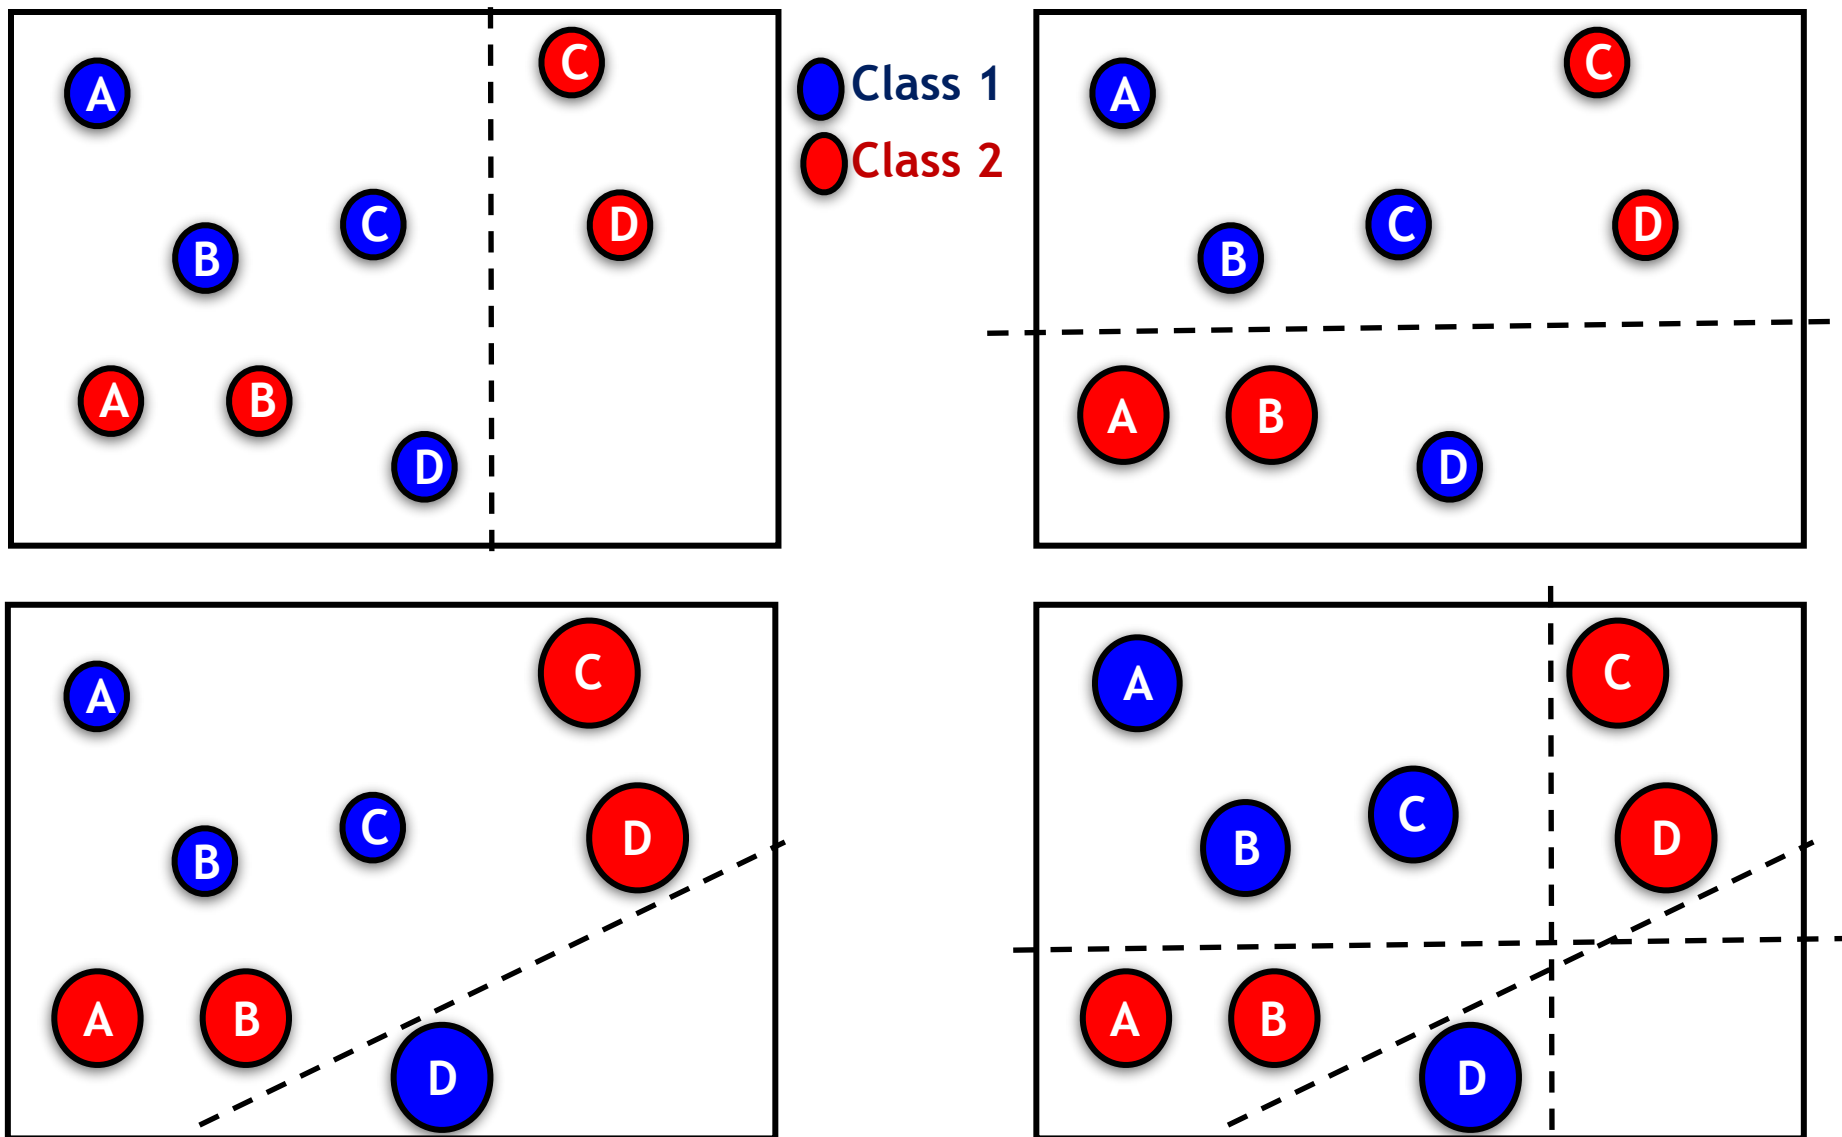
For a new testing point  $(x', y')$ ,

$$H(x') = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x') \right)$$

---

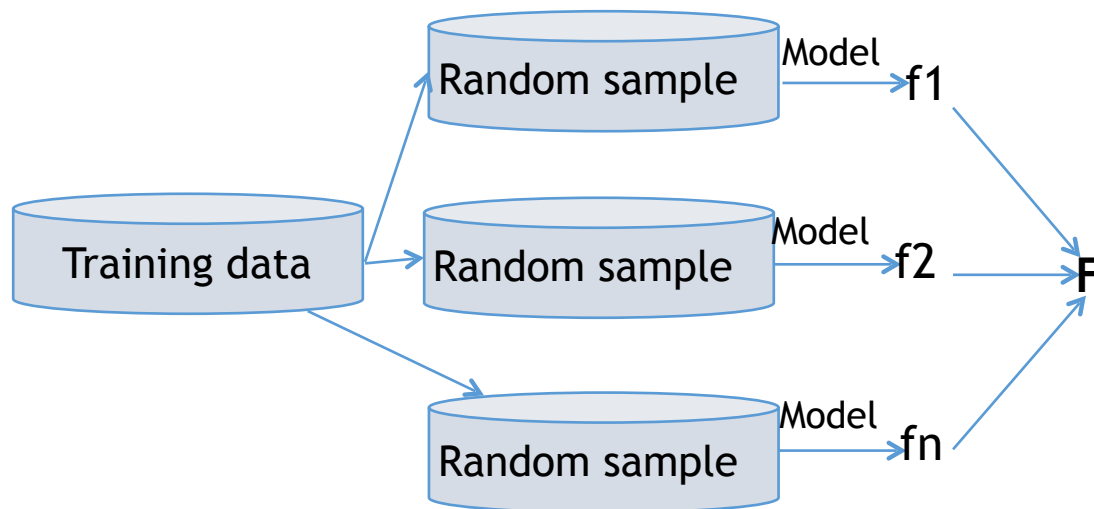


# Boosting illustration



# Bagging VS Boosting

## Bagging

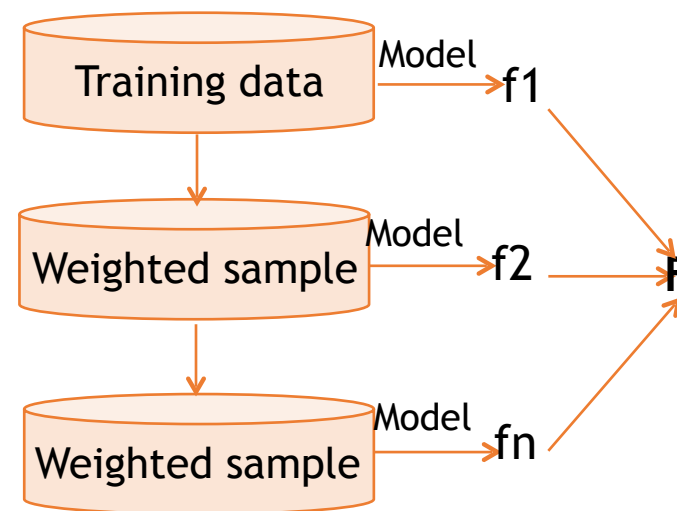


Resampling

Uniform distribution

Parallel Style

## Boosting



Reweighting

Non-uniform distribution

Sequential Style

# Gradient Boosted Models (GBM's)

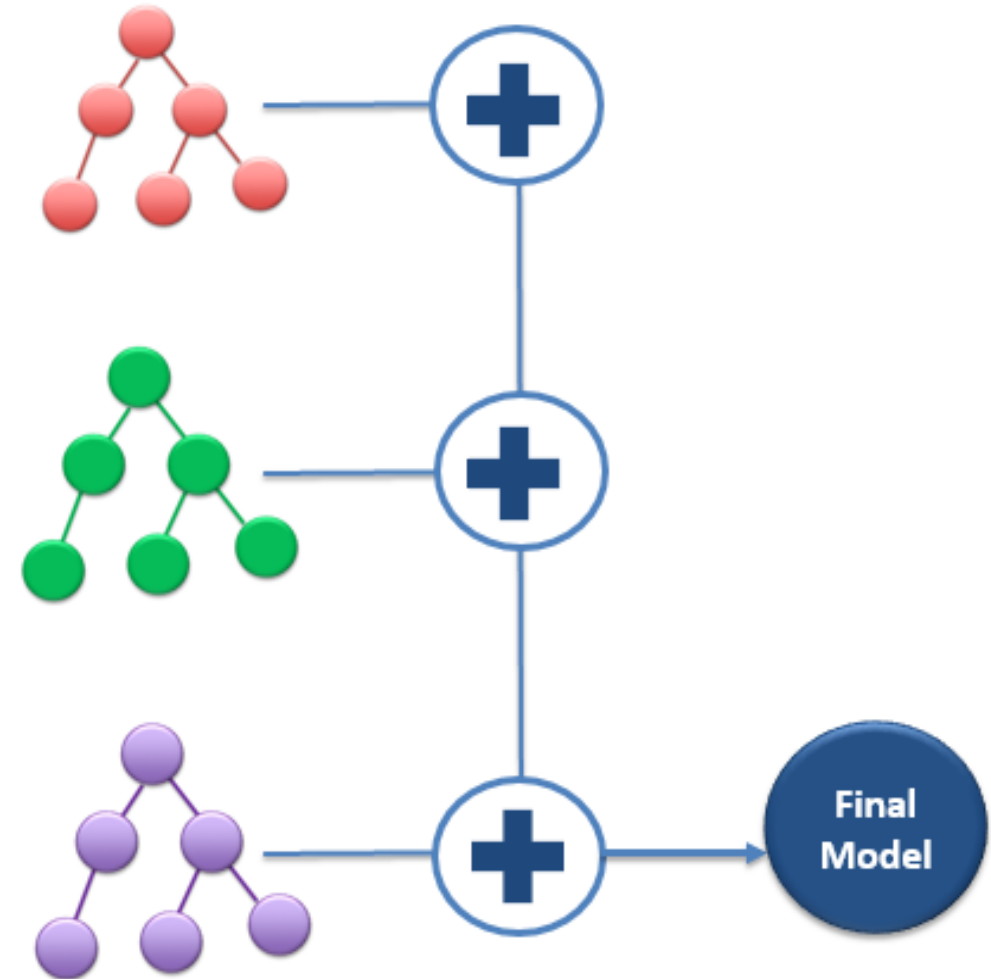
- In gradient boosting, it trains many model sequentially. Each new model gradually minimizes the loss function using Gradient Descent method.
- The learning procedure consecutively fit new models to provide a more accurate estimation of the response variable.
- *In Adaboost, the weights are derived from the misclassifications of the previous model. The resulting increased weights assigned to misclassifications.*
- The result of Gradient Boosting is an altogether different function from the beginning, because the result is the addition of multiple functions.

# Gradient Boosting Trees

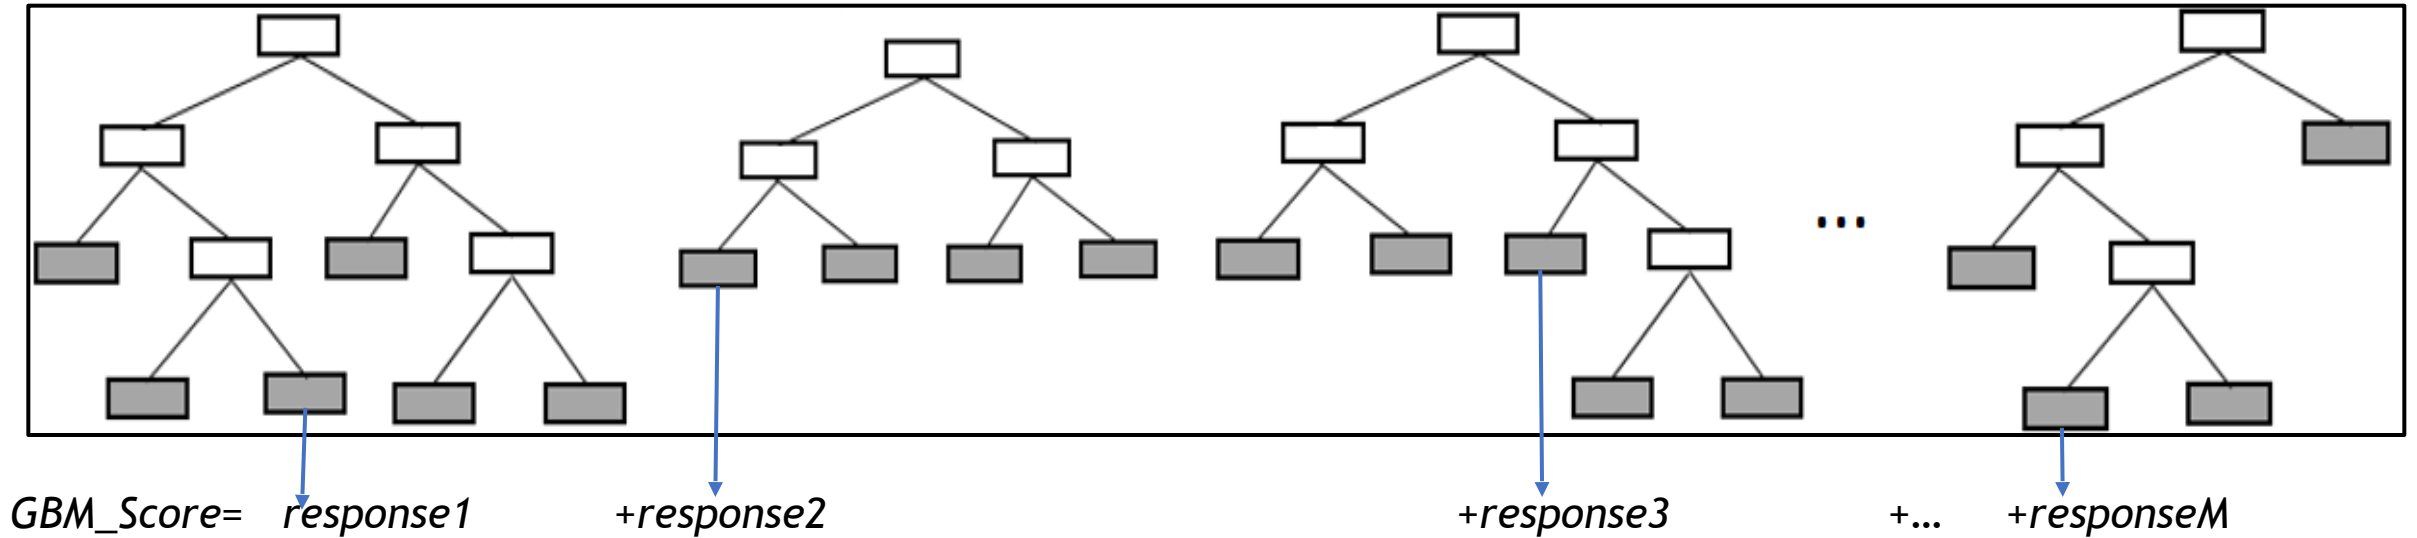
At each iteration:

- Draws a subsample of the training data (**without replacement**)
- Constructs a regression tree from the sample
- All trees are added together to get the final model

*Jerome H. Friedman, "Stochastic gradient boosting", Computational Statistics & Data Analysis 2002*

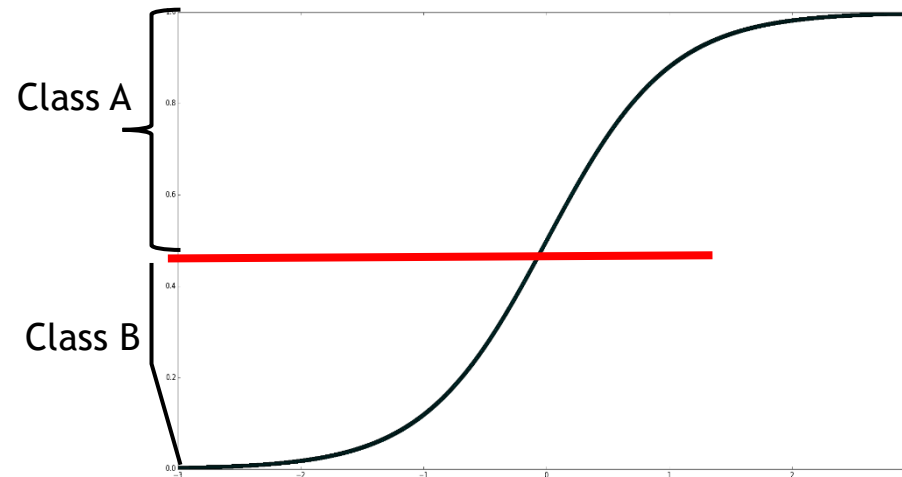


# Prediction with GBM



$$Pred = \frac{e^{GBM\_score}}{e^{GBM\_score} + e^{-GBM\_score}}$$

*monotonic function*



# AdaBoost & GBM in SKlearn

`sklearn.ensemble.AdaBoostClassifier`

Input parameters:

`base_estimator`  
`n_estimators`  
`Learning_rate`  
...

`sklearn.ensemble.GradientBoostingClassifier`

Input parameters:

`loss`  
`n_estimators`  
`learning_rate`  
`n_estimator`  
`max_depth`  
`Max_features`  
...

# XGBoost

- XGBoost is an implementation of GBM, with major improvements.
- GBM's build trees sequentially, but XGBoost is parallelized. This makes XGBoost faster.
- XGBoost is an open-sourced machine learning library available in Python, R, Julia, Java, C++, Scala.

# XGBoost features

## 1. Split finding algorithms: approximate algorithm

- Candidate split points are proposed based on the percentiles of feature distribution.
- The continuous features are binned into buckets that are split based on the candidate split points.
- The best solution for candidate split points is chosen from the aggregated statistics on the buckets.

## 2. Column block for parallel learning

- To reduce sorting costs, data is stored in in-memory units called 'blocks'.
- Each block has data columns sorted by the corresponding feature value.
- This computation needs to be done only once before training and can be reused later.
- Sorting of blocks can be done independently and divided between parallel threads.
- The split finding can be parallelized as the collection of statistics for each column is done in parallel.



# XGBoost features ..cont.

## **3. Sparsity-aware algorithm:**

- XGBoost visits only the default direction (non-missing entries) in each node.

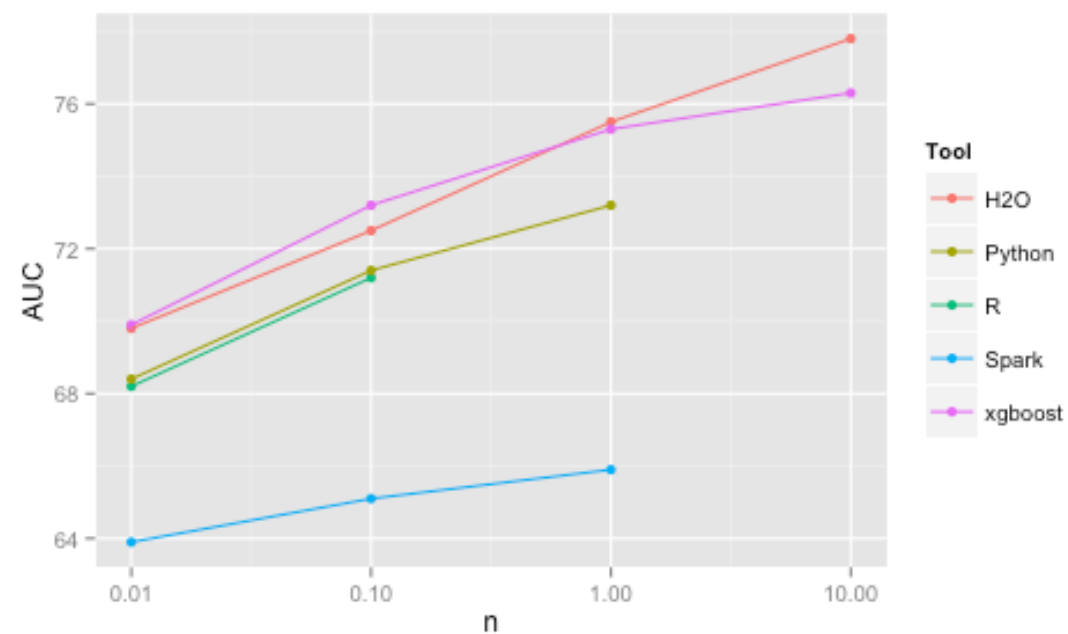
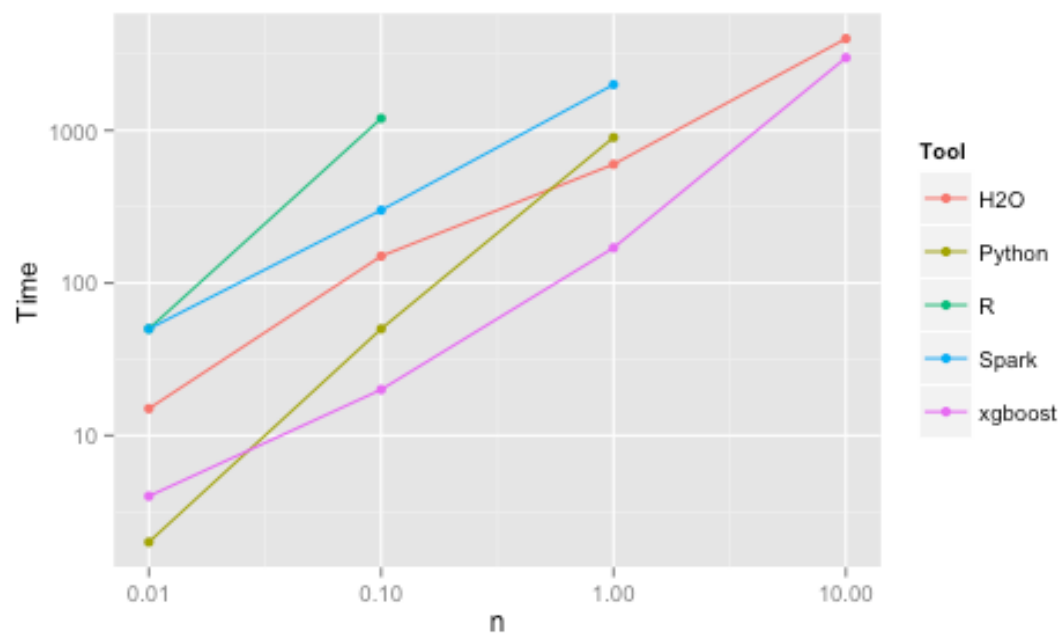
## **4. Cache-aware access:**

- Optimizes how many examples per block.

## **5. Out-of-core computation:**

- For blocks that does not fit into memory, they are compressed on the disk .
- The blocks are decompressed on the fly (In parallel).

# Xgboost VS Other tools



<http://datascience.la/benchmarking-random-forest-implementations/>

# Light GBM...the latest of Boosting algorithms

- A fast, distributed, high performance gradient boosting framework.
- Used for ranking, classification and many other machine learning tasks.
- It is under the umbrella of the DMTK(<http://github.com/microsoft/dmtk>) project of Microsoft.
- Similar to XGBoost but it is faster

# Advantages Ensemble Learning

Accuracy

*less*  
Variance

*less*  
Overfitting

Diversity

In Ensemble learning, the large variance of unstable learners is 'averaged out' across multiple learners.

Different classifiers to work on different random subsets of the full feature space or different subset of the training data.

Imagine we have:

- An ensemble of 5 independent classifiers.
- Accuracy is 70% for each

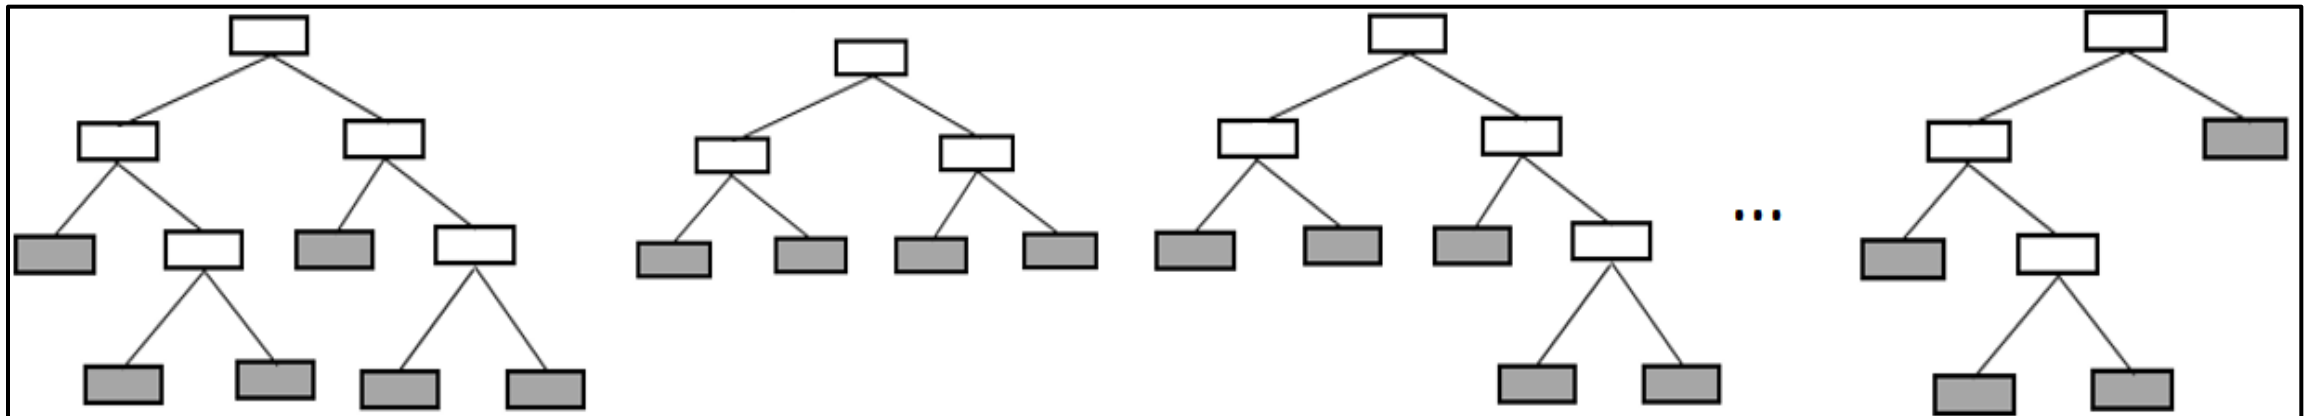
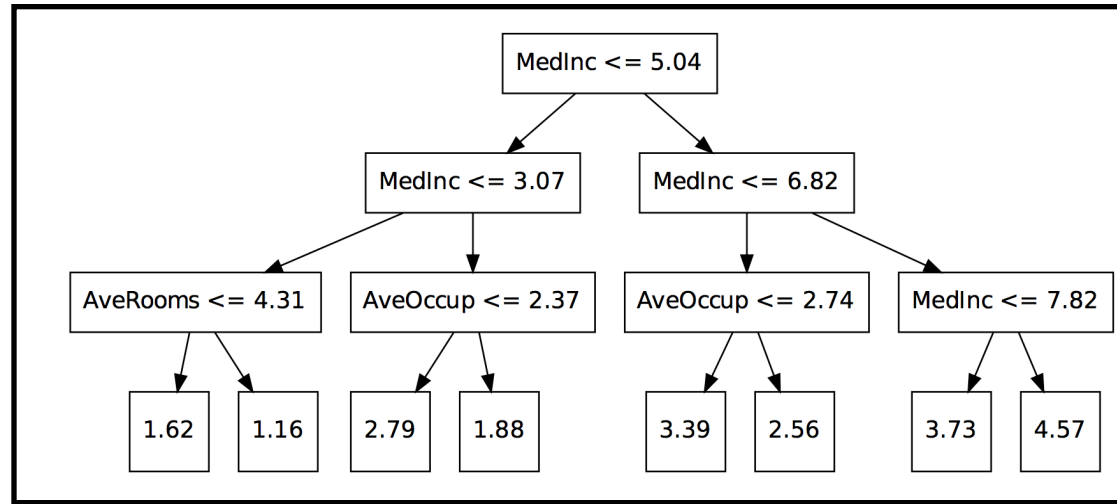
What is the accuracy for the majority vote?

$$10 (.7^3)(.3^2)+5(.7^4)(.3)+(.7^5)$$

**83.7% majority vote accuracy**

How about if we have 101 such classifiers  
**99.9% majority vote accuracy**

# Are ensembles easy to understand?



# Are ensembles easy to understand?

- Decision trees are easy to understand
- Ensemble models are considered complex model, interpreting predictions from them is a challenge.

## **LIME - Local Interpretable Model-Agnostic Explanations**

Approximate the complex model near a given prediction. - Ribeiro et al. 2016

## **SHAP (SHapley Additive exPlanations)**

A unified approach to interpreting model predictions, Scott Lundberg, Su-In Lee 2017



# References

- Zhou, Zhi-Hua, “Ensemble Learning”, Encyclopedia of Biometrics, 2009
- *Dietterich, T. Ensemble methods in machine learning. Multiple classifier systems, Vol. 1857 of Lecture Notes in Computer Science. 2000*
- Ho, Tin Kam, The Random Subspace Method for Constructing Decision Forests. IEEE Transactions on Pattern Analysis and Machine Intelligence, 1998.
- VALENTINI, Giorgio, and Francesco MASULLI, 2002. Ensembles of Learning Machines. 2002. Revised Papers, Volume 2486 of Lecture Notes in Computer Science. Berlin: Springer, pp. 3-19.
- FREUND, Yoav, and Robert E. SCHAPIRE, 1996. Experiments with a New Boosting Algorithm. In: Lorenza SAITTA, ed. (ICML '96). San Francisco, CA: Morgan Kaufmann, pp. 148-156.
- SCHAPIRE, Robert E., 1990. The Strength of Weak Learnability. Machine Learning, 5(2), 197-227. Oxford University Press.
- BREIMAN, Leo, 1996. Bagging Predictors. Machine Learning, 24(2), 123-140. Sewell M (2011) Ensemble learning, Technical Report RN/11/02. Department of Computer Science, UCL, London
- G Brown, “Ensemble Learning”, Encyclopedia of Machine Learning, 2010
- <http://scikit-learn.org/0.16/modules/ensemble.html>
- <http://scikit-learn.org/stable/modules/multiclass.html>
- <http://xgboost.readthedocs.io/en/latest/>



# Boosting Algorithms

Omar Odibat, Data Scientist, **VISA**



**Thank you!!!**